

COSC329 project report

Zhenghua Chen

My project is to analyse the required skills for Data Scientist and Software Developer in Canada.

Scraping

1. prepare

Before scraping, I firstly looked at the indeed website in Canada <https://ca.indeed.com> to get a general idea. After I searched in the box, the web page will show pages of results and each result has a link which leads to that specific job page.

My general method is to go into the search page first, get all the result links on that page, read each job's description through each link and then go to the next page, do the same thing until I collect enough posts.

2. pipeline

I used python package requests and BeautifulSoup for completing this pipeline. All the links in the search page were identified by looking for all the a (link) item under the div item with id "mosaic-provider-jobcards". After a posting link is accessed, its job description was identified by the div item with id: "jobDescriptionText" and class "jobsearch-jobDescriptionText". The next page on the search page was found by ul item with class "pagination-list" and the link item with "aria-label":"Next" inside the class.

Some links lead to the company's own website which don't have the exact class as above, but it's only a few posts. In that case I output a link of this job and the job description can be manually extracted later.

I make the process sleep for 3 seconds before scraping each page to avoid high frequency. A random user agent from a list was used in requests header to keep from being blocked by the website.

After I scraped one job posting, I wrote it directly to a csv row with job title, location, description and link for later use.

3. results

Though the code can finish successfully with no error, the jobs I got were far less from the number of search results showed on the website. I only got 1114 Software Developer job postings and 939 Data Scientist job postings from indeed, while the total jobs would be over 10000. The scraped jobs were also not in the same order as I saw on the web page. I checked the csv file and it looked good with correct information scraped. I didn't figure out why this happened so I left this as a unfinished feature of my project.

Analysis

1. prepare

Before processing the csv file with all the job information, I searched online for effective methods. The most popular methods for extracting text features online were bag-of-words and tfidf vectorization. I adopted the python package `sklearn.feature_extraction.text` and the useful tools `CountVectorizer` and `TfidfVectorizer` to try on my work. I also decided to use the package `RegEx` for cleaning and `nltk.wordpunct_tokenize` for tokenizing the strings. I used the `lemmatize` function from the package `textblob` for lemmatization and English stop words from `nltk.corpus` for removing unrelated words. `Pandas.DataFrame` and `numpy` were also utilized for dealing with large amount of data.

2. methods

The first step overall was to read the data from csv file and organize them in the pandas dataframe. I removed the rows with NaN and then did the general cleaning by removing special characters, digits and underscores. Then all the texts were converted to lowercase.

I developed two methods for extracting features. The first one is to use `tfidf-vectorizer` from the package directly.

In this method I first cleaned data again by removing non-english words, removing stop words, lemmatization and dropping words less than 2 chars. When I was developing and testing on the workflow, I noticed that there were still a lot of noise in the text, so I added an input of list which declares the extra stop words that need to be removed. Then the cleaned text string was converted to word list and was put into the `TfidfVectorizer` with `max_df=2` and `min_df=0.5`. After running the function `fit_transform` on the corpus, all the features were extracted and

their tfidf scores were listed for each job. Then I get the top n features by averaging the tfidf scores for each feature and ranking them.

The second method was more of my own techniques. I found out that only using tfidf vectorizer and basic cleaning did not work very well since there were still a large number of less-important features about the company, general description or other noise. To get a more precise outcome, I believed that the texts need to be further preprocessed. Based on our searching habits, I created a customized searching method to look for the section in the job posting named “requirements”, “qualifications” or “skills”. Though a lot of jobs do not have these section names for skills but have “who you are”, “what you’ll need to do”, around 5/7 of all the posts contained these noticeable keywords. Thus, I tried to extract only a proportion of words following these key words. Notice that because of web scraping, sometimes these keywords were concatenated with other words such as “a sentence endSkills you need....”. In this way we should search for the substring instead of the tokenized words.

Since I didn’t want to clean too much to before finding the patterns for this approach, I just removed the stop words and lemmatized the text so that “skills” would all become “skill” for easier detection. Then I extracted 80 words after each occurrence of the keywords ["requirement", "qualification", "skill"] for each text. If none of the keywords appeared, I searched for the keyword substrings in the text string and extracted the following 640 chars. Finally, the extra stop words were removed and a list of words were gained from the extracted string. Then the tfidf and count vectorizers were applied on these words.

Finally the top 20 features of each method were showed and I calculated the overlap of them using `bhattacharyya_coefficient` function from `distance` package. A heatmap was plotted to show the final comparison.

3. tests and results

Data Scientist:

For method 1:

```
[('data', 0.4745309723037524), ('experience', 0.2562409372411891), ('learn', 0.21741967975425838), ('business', 0.19430043251834614), ('team', 0.18543229194463734), ('machine', 0.1529591042978648), ('machine learn', 0.14845198480721547), ('analytics', 0.13411077451607425), ('science', 0.12741647448115886), ('build', 0.11699428945000243), ('analysis', 0.1059335505237994), ('engineer', 0.09744002242886084), ('new', 0.09501028884584822), ('knowledge', 0.09391096055039992), ('data science', 0.09095950302422869), ('help', 0.08879748154398966), ('design', 0.08771769283194432), ('support', 0.08477634817727388), ('technical', 0.08282301184291198), ('understand', 0.08174505460120045)]
```

For method 2:

Count vectorization:

```
[('data', 7.4185303514377), ('experience', 5.567625133120341), ('business', 2.617678381256656), ('learn', 2.5228966986155483), ('team', 2.252396166134185), ('science', 1.8519701810436635), ('machine', 1.516506922257721), ('machine learn', 1.4494142705005324), ('model', 1.428115015974441), ('analytics', 1.3258785942492013), ('knowledge', 1.3194888178913737), ('analysis', 1.2523961661341854), ('python', 1.1448349307774228), ('process', 1.106496272630458), ('engineer', 1.066027689030884), ('problem', 1.0425985090521832), ('degree', 1.012779552715655), ('project', 1.0021299254526093), ('understand', 0.9957401490947817), ('data science', 0.9850905218317358)]
```

TFIDF vectorization:

```
[('data', 0.2877633101224421), ('experience', 0.2756532150624001), ('learn', 0.15325665146712186), ('business', 0.14650920636848752), ('team', 0.1442797654502848), ('science', 0.1101214687427899), ('knowledge', 0.10052509021414582), ('machine', 0.09537498590419481), ('model', 0.09303194145428996), ('machine learn', 0.09281667674616376), ('analytics', 0.08670057399632608), ('python', 0.08591469951180035), ('process', 0.0827930727525265), ('engineer', 0.08192012619924995), ('analysis', 0.07932254545507077), ('problem', 0.07841315525590559), ('project', 0.07561435962042432), ('understand', 0.07076068422663422), ('degree', 0.06712962710563769), ('data science', 0.06681374071601766)]
```

From results above, we can see the most important feature is data, which is apparent for this job. Experience is also very important for any industry work in common sense. It also contains “business”, “machine learn”, “model”, “analytics” and “python”, which are all must-have skills to be a data scientist. We can also conclude that most of the jobs of data scientist are related to business company. It is obvious since most business companies need to process and analyse the huge amount of data behind customer behaviors to improve their service correspondingly. Machine learning is especially popular among all the techniques due to its power of identifying, predicting and extracting. A degree is desired for this position too.

Software Developer:

For method 1:

```
[('experience', 0.36315990797403175), ('team', 0.2899107897934766), ('design', 0.1989588877278291), ('test', 0.16643072740503162), ('code', 0.1604026579337243), ('knowledge', 0.15305727672595665), ('build', 0.14914541087724317), ('technical', 0.13604905620153687), ('technology', 0.13468307185121287), ('environment', 0.1340344319246207), ('er', 0.13389282520859083), ('computer', 0.13382761218268285), ('company', 0.13219381676187777), ('job', 0.125825006138966), ('full', 0.12391009289097059), ('look', 0.1135175757275613)]
```

For method 2:

Count vectorization:

```
[('experience', 4.693140794223827), ('software', 3.1958483754512637), ('design', 2.1796028880866425), ('team', 2.032490974729242), ('test', 1.78971119133574), ('application', 1.6525270758122743), ('knowledge', 1.48014440433213), ('technology', 1.3005415162454874), ('system', 1.2851985559566788), ('technical', 1.2382671480144405), ('code', 1.227436823104693), ('computer', 1.167870036101083), ('programm', 0.990072202166065), ('understand', 0.9575812274368231), ('environment', 0.9467509025270758), ('engineer', 0.9142599277978339), ('web', 0.9106498194945848), ('process', 0.842057761732852), ('communication', 0.8312274368231047), ('degree', 0.8185920577617328)]
```

TFIDF vectorization:

```
[('experience', 0.2790188128574318), ('software', 0.1771664861509763), ('team', 0.14653905103455686), ('design', 0.12570037361211184), ('knowledge', 0.11043349699422986), ('test', 0.11042675687065157), ('application', 0.10563755812287254), ('technology', 0.1031428028807942), ('code', 0.09264182875744362), ('technical', 0.08866975924485232), ('computer', 0.0857308615819803), ('environment', 0.08293039167910356), ('understand', 0.08191492583501561), ('web', 0.08173344515569282), ('system', 0.08120392553333655), ('programm', 0.07673117613265693), ('engineer', 0.07405795043957615), ('communication', 0.0728312294025619), ('degree', 0.07221857784491101), ('process', 0.06680401883645198)]
```

From the above results about software developer, we can see that the most important feature is experience instead of data. That's the biggest difference from data scientist jobs since developers don't need to deal with a large size of data. Their jobs are mostly about: "software", "design", "test", "application", "code", "web", "system" and "program". So, it's not surprising to see that these are the required skills.

To compare the preprocessing techniques of the two methods, I further checked the extracted text before vectorization of a random data scientist job to see the difference.

Data Scientist checking...

After cleaning using method 1:

challenge company organization role collaborate platform global collaborative present problem goal speed knowledge capital technology scientist help equity python computer page job support keep growth complex structured every acumen better passion think analytical make per combine degree everyone aggregate paced create raw intelligence field active analyze presentation research industry eye extract business algebra data preferred learn full team knack private engineer across information innovative rapidly revolution

ionize culture based compensation math aptitude predictive experience ecosystem collection tech undertake vision large discover competitive proprietary graduate type understand relevant communication innovation highly identify opportunity mind modern well day statistic access rely backed proven north efficiency machine propose ensemble earn funded neural build quantitative visualization tableau market globally excellent science environment mean unique fast valuable grow look

After cleaning using method 2:

ba challenge company organization global tensorflow us collaborative keras solve problem knowledge capital scientist equity python computer job support growth complex every leadership come acumen better think analytical degree evolve paced intelligence field active presentation dealmaker industry algebra scikit business data preferred learn sql team engineer innovative rapidly revolutionize culture based compensation math aptitude experience bachelor champion competitive graduate understand relevant communication model opportunity frame nlp mind well modern day statistic backed proven net machine earn funded neural quantitative tableau bert u market excellent science environment g canadian fast grow

We can see that the first preprocessed text had more general information like “collaborate”, “analytical”, “communication” which are required by most of the industries. The second cleaned text, however, revealed more specific skills that a data scientist should have, such as “tensorflow”, “keras”, “scikit”, “sql”, “nlp”, “bert”... They both contained “python”, “math”, “computer”, so these can be considered as the fundamental skills for both jobs.

Then the Bhattacharyya coefficient was calculated as 0.8261 which indicated a large overlap between the two jobs. This makes sense since they are all computer science related jobs and have a lot of required skills in common.

The following heatmap further explores the overlapping features between the two jobs. The most important feature of software developer jobs, “experience”, is also the second important feature of data scientist jobs. This together with other 6 common features result in a high Bhattacharyya coefficient.

